

# Data Analysis and Visualization Exercise

## 2.2: Data import

*Felix Brechtmann, Matthias Heinig, Vicente Yépez*

30 October 2019

### 1 Flat file questions

---

#### 1.1 Question

Read the titanic file (titanic.csv) and figure out who was the oldest surviving passenger of the titanic accident? Hint: `?subset`

```
tit_df <- read.csv("./extdata/titanic.csv")
head(tit_df)
##   pclass survived                name    sex
## 1      1         1 Allen, Miss. Elisabeth Walton female
## 2      1         1 Allison, Master. Hudson Trevor    male
## 3      1         0 Allison, Miss. Helen Loraine female
## 4      1         0 Allison, Mr. Hudson Joshua Creighton    male
## 5      1         0 Allison, Mrs. Hudson J C (Bessie Waldo Daniels) female
## 6      1         1 Anderson, Mr. Harry    male
##   age sibsp parch ticket    fare  cabin embarked boat body
## 1 29.00     0     0 24160 211.3375    B5      S      2   NA
## 2  0.92     1     2 113781 151.5500 C22 C26      S     11   NA
## 3  2.00     1     2 113781 151.5500 C22 C26      S      NA
## 4 30.00     1     2 113781 151.5500 C22 C26      S     135
## 5 25.00     1     2 113781 151.5500 C22 C26      S      NA
## 6 48.00     0     0 19952 26.5500  E12      S      3   NA
##               home.dest
## 1                St Louis, MO
## 2 Montreal, PQ / Chesterville, ON
## 3 Montreal, PQ / Chesterville, ON
## 4 Montreal, PQ / Chesterville, ON
## 5 Montreal, PQ / Chesterville, ON
## 6                New York, NY
# restrict to survivors and get oldest; show name
subset( tit_df, survived==1 & age == max(age, na.rm=T), name)
##               name
## 15 Barkworth, Mr. Algernon Henry Wilson
```

## Data Analysis and Visualization Exercise 2.2: Data import

### 1.2 Question

For the next two questions we simulate files as strings. They can be read, as if they were files.

A csv file has numbers as column names in the first row. Which parameter of `read.table()` needs to be adjusted to read the column names as they are in the csv?

```
tmp_tidy_table <- "1_colname,2_colname,3_colname
3,4,5
a,b,c"
read.csv(text = tmp_tidy_table)
##   X1_colname X2_colname X3_colname
## 1         3         4         5
## 2         a         b         c
```

```
# parameter `check.names`: a logical, tests for syntactically valid variable names
tidy_txt_df <- read.csv(text = tmp_tidy_table, check.names = FALSE)
tidy_txt_df
##   1_colname 2_colname 3_colname
## 1         3         4         5
## 2         a         b         c
```

### 1.3 Question

How to read the following table to have the `identical()` information as in `tidy_txt_df` from question above?

```
tmp_messy_table <- "# This line is just useless info

1_colname,2_colname,3_colname
3,4,5

a,b,c"
```

```
# parameter `comment.char`: a character vector of length one
#   containing a single character or an empty string.
#   Use "" to turn off the interpretation of comments altogether.
# The parameter `blank.lines.skip` is TRUE by default.
messy_txt_df <- read.csv(text = tmp_messy_table, check.names = F, comment.char = '#')
identical(tidy_txt_df, messy_txt_df)
## [1] TRUE
```

## 2 Excel questions

### 2.1 Question

Read only Name, Type and Total columns for only the first 10 pokemons of the pokemon.xlsx file. Hint: take a look at the file using Excel or any other spreadsheet application.

```
library(readxl)
poke_file <- file.path('extdata/pokemon.xlsx')
poke_df <- read_excel(poke_file, sheet='Pokemon', range="B1:D11")
poke_df
## # A tibble: 10 x 3
##   Name      Type  Total
##   <chr>    <chr> <dbl>
## 1 Bulbasaur GRASS   318
## 2 Bulbasaur POISON  318
## 3 Ivysaur   GRASS   405
## 4 Ivysaur   POISON  405
## 5 Venusaur  GRASS   525
## 6 Venusaur  POISON  525
## 7 Mega Venusaur GRASS  625
## 8 Mega Venusaur POISON  625
## 9 Charmander FIRE    309
## 10 Charmeleon FIRE    405
```

### 2.2 Question

Using the summer\_olympic\_medals.xlsx file, which athlete won most bronze medals?

```
oly_file <- file.path('extdata/summer_olympic_medals.xlsx')
```

```
library(readxl)
oly_df <- read_excel(oly_file, sheet='ALL MEDALISTS')
head(oly_df)
## # A tibble: 6 x 10
##   City  Edition Sport Discipline Athlete NOC   Gender Event Event_gender
##   <chr>   <dbl> <chr> <chr>      <chr> <chr> <chr> <chr> <chr>
## 1 Athe~  1896 Aqua~ Swimming HAJOS,~ HUN   Men   100m~ M
## 2 Athe~  1896 Aqua~ Swimming HERSCH~ AUT   Men   100m~ M
## 3 Athe~  1896 Aqua~ Swimming DRIVAS~ GRE   Men   100m~ M
## 4 Athe~  1896 Aqua~ Swimming MALOKI~ GRE   Men   100m~ M
## 5 Athe~  1896 Aqua~ Swimming CHASAP~ GRE   Men   100m~ M
## 6 Athe~  1896 Aqua~ Swimming CHOROP~ GRE   Men   1200~ M
## # ... with 1 more variable: Medal <chr>

# There are different solutions for this
oly_dt <- as.data.table(oly_df)
bronze <- oly_dt[Medal == "Bronze",]
```

## Data Analysis and Visualization Exercise 2.2: Data import

```
# 1. Using table() function
bronze_counts <- table(subset(oly_dt, Medal == "Bronze", "Athlete"))
head(sort(bronze_counts, decreasing = T))
##
##              NEMOV, Alexei              OTTEY-PAGE, Merlene
##              6              6
##              SAVOLAINEN, Heikki              VAN ALMSICK, Franziska
##              6              6
##              BUSCHSCHULTE, Antje DE JONG, Adrianus Egbertus Willem
##              5              5
bronze_athl <- bronze_counts[ bronze_counts==max(bronze_counts)]
bronze_athl
##
##              NEMOV, Alexei              OTTEY-PAGE, Merlene              SAVOLAINEN, Heikki
##              6              6              6
## VAN ALMSICK, Franziska
##              6
# View their record
# subset(oly_df, Athlete %in% names(bronze_athl),
#       c('Athlete', 'Edition', "Sport", "Event", "Medal"))

# 2.Using aggregate() function
bronze.l <- as.data.table(aggregate(bronze,
                                   by=list(bronze$Athlete), FUN=length))
bronze.l[Medal == max(Medal),]
##              Group.1 City Edition Sport Discipline Athlete NOC Gender
## 1:              NEMOV, Alexei      6      6      6      6      6      6      6
## 2:              OTTEY-PAGE, Merlene      6      6      6      6      6      6      6
## 3:              SAVOLAINEN, Heikki      6      6      6      6      6      6      6
## 4: VAN ALMSICK, Franziska      6      6      6      6      6      6      6
##      Event Event_gender Medal
## 1:      6      6      6
## 2:      6      6      6
## 3:      6      6      6
## 4:      6      6      6

# 3. Using dcast() function. More on this in Tidy Data lecture
ox <- dcast(oly_dt, Athlete ~ Medal)
ox <- subset(ox, Bronze == max(Bronze))

# 4. Using .N command from data.table. More to this in Data Table lecture
bronze[, N := .N, by = Athlete]
bronze[N == max(N), unique(Athlete)]
## [1] "SAVOLAINEN, Heikki" "OTTEY-PAGE, Merlene"
## [3] "VAN ALMSICK, Franziska" "NEMOV, Alexei"
```

## Data Analysis and Visualization Exercise 2.2: Data import

### 2.3 Question

Are the columns `Gender` and `Event_gender` consistent? Find inconsistent gender entries.

```
# There was a male Bronze-medal winner in ladies marathon in 2000.
unique(oly_df$Gender)
## [1] "Men" "Women"
unique(oly_df$Event_gender)
## [1] "M" "X" "W"
subset(oly_df, Gender=='Men' & Event_gender=='W')
## # A tibble: 1 x 10
##   City Edition Sport Discipline Athlete NOC Gender Event Event_gender
##   <chr>   <dbl> <chr> <chr>      <chr>  <chr> <chr> <chr> <chr>
## 1 Syd~   2000 Athl~ Athletics CHEPCH~ KEN Men   mara~ W
## # ... with 1 more variable: Medal <chr>
subset(oly_df, Gender=='Women' & Event_gender=='M')
## # A tibble: 0 x 10
## # ... with 10 variables: City <chr>, Edition <dbl>, Sport <chr>,
## #   Discipline <chr>, Athlete <chr>, NOC <chr>, Gender <chr>, Event <chr>,
## #   Event_gender <chr>, Medal <chr>
```

### 2.4 Question

Which country won most medals? Which country has the highest ratio of silver medals? Use the data in the country summary sheet starting at row 147 of the `summer_olympic_medals.xlsx` file.

```
# There is also a summary sheet for nations
nation_medal_df <- read_excel(oly_file, sheet='COUNTRY TOTALS', range="A147:F286")

head(nation_medal_df)
## # A tibble: 6 x 6
##   NOC      Country      Bronze Gold Silver `Grand Total`
##   <chr>   <chr>      <dbl> <dbl> <dbl>      <dbl>
## 1 Grand Total <NA>      9689  9850  9677      29216
## 2 USA      United States  1052  2088  1195      4335
## 3 URS      Soviet Union    584   838   627      2049
## 4 GBR      United Kingdom  505   498   591      1594
## 5 FRA      France         475   378   461      1314
## 6 GER      Germany        454   407   350      1211
# Remove Grand.Total row
nation_medal_df <- subset(nation_medal_df, !is.na(Country))
# Get max
subset(nation_medal_df, `Grand Total`==max(`Grand Total`))
## # A tibble: 1 x 6
##   NOC      Country      Bronze Gold Silver `Grand Total`
##   <chr> <chr>      <dbl> <dbl> <dbl>      <dbl>
## 1 USA      United States  1052  2088  1195      4335
```

## Data Analysis and Visualization Exercise 2.2: Data import

```
nation_medal_df[, 'silver_ratio'] <- with(nation_medal_df, Silver/`Grand Total`)
# alternatively:
# nation_medal_df[, 'silver_ratio'] <- nation_medal_df$Silver/nation_medal_df$`Grand Total`
subset(nation_medal_df, silver_ratio==max(silver_ratio, na.rm=T))
## # A tibble: 13 x 7
##   NOC   Country      Bronze   Gold Silver `Grand Total` silver_ratio
##   <chr> <chr>      <dbl>   <dbl> <dbl>      <dbl>      <dbl>
## 1 PAR   Paraguay      NA     NA    17         17         1
## 2 SCG   Serbia        NA     NA    14         14         1
## 3 NAM   Namibia        NA     NA     4          4         1
## 4 SIN   Singapore      NA     NA     4          4         1
## 5 SRI   Sri Lanka      NA     NA     2          2         1
## 6 TAN   Tanzania       NA     NA     2          2         1
## 7 VIE   Vietnam        NA     NA     2          2         1
## 8 AHO   Netherlands Antilles* NA     NA     1          1         1
## 9 CIV   Cote d'Ivoire   NA     NA     1          1         1
## 10 ISV  Virgin Islands* NA     NA     1          1         1
## 11 SEN   Senegal        NA     NA     1          1         1
## 12 SUD   Sudan          NA     NA     1          1         1
## 13 TGA   Tonga          NA     NA     1          1         1
```

## 2.5 Question

Which countries did participate, but without winning medals? Assume, that all countries listed in the IOC COUNTRY CODES sheet participated.

```
participants <- read_excel(oly_file, sheet='IOC COUNTRY CODES', range="A1:C202")
# workaround
# participants <- read.table(
#   sub('.xlsx', '_IOC_COUNTRY_CODES.csv', oly_file),
#   sep=':', quote='@', header=T
# )
head(participants)
## # A tibble: 6 x 3
##   Country      `Int Olympic Committee code` `ISO code`
##   <chr>      <chr>                        <chr>
## 1 Afghanistan AFG                                AF
## 2 Albania     ALB                                AL
## 3 Algeria     ALG                                DZ
## 4 American Samoa* ASA                                AS
## 5 Andorra     AND                                AD
## 6 Angola      ANG                                AO

## make sure to have proper variable names
colnames(participants) <- make.names(colnames(participants))

no_medals <- setdiff(participants$Int.Olympic.Committee.code, nation_medal_df$NOC)
length(no_medals)
## [1] 78
```

## Data Analysis and Visualization Exercise 2.2: Data import

```
c(subset(participants, Int.Olympic.Committee.code %in% no_medals, "Country"))
## $Country
## [1] "Albania" "American Samoa*"
## [3] "Andorra" "Angola"
## [5] "Antigua and Barbuda" "Aruba*"
## [7] "Bahrain" "Bangladesh"
## [9] "Belize" "Benin"
## [11] "Bhutan" "Bolivia"
## [13] "Bosnia and Herzegovina" "Botswana"
## [15] "British Virgin Islands" "Brunei"
## [17] "Burkina Faso" "Cambodia"
## [19] "Cape Verde" "Cayman Islands*"
## [21] "Central African Republic" "Chad"
## [23] "Comoros" "Congo"
## [25] "Congo, Dem Rep" "Cook Islands"
## [27] "Cyprus" "Dominica"
## [29] "East Timor (Timor-Leste)" "El Salvador"
## [31] "Equatorial Guinea" "Fiji"
## [33] "Gabon" "Gambia"
## [35] "Grenada" "Guam"
## [37] "Guatemala" "Guinea"
## [39] "Guinea-Bissau" "Honduras"
## [41] "Jordan" "Laos"
## [43] "Lesotho" "Liberia"
## [45] "Libya" "Liechtenstein"
## [47] "Madagascar" "Malawi"
## [49] "Maldives" "Mali"
## [51] "Malta" "Mauritania"
## [53] "Micronesia" "Monaco"
## [55] "Burma" "Nauru"
## [57] "Nepal" "Nicaragua"
## [59] "Oman" "Palau"
## [61] "Palestine, Occupied Territories" "Papua New Guinea"
## [63] "Romania" "Rwanda"
## [65] "Saint Kitts and Nevis" "Saint Lucia"
## [67] "Saint Vincent and the Grenadines" "Samoa"
## [69] "San Marino" "Sao Tome and Principe"
## [71] "Seychelles" "Sierra Leone"
## [73] "Solomon Islands" "Somalia"
## [75] "Swaziland" "Turkmenistan"
## [77] "Vanuatu" "Yemen"
```

## 3 SQL questions

### 3.1 Question

Connect to the `extdata/Northwind.sl3` SQLite data base (using the 'RSQLite' package). Inspect the data base tables using the 'dbListTables' and 'dbListFields' functions. Put together a SQL statement to retrieve a table that lists for all customers (name of the company, name of the contact person and city) all the products (name of the product) that they ordered. Execute the statement using 'dbGetQuery'. How many rows does this table have? Display the first 5 rows.

We provide the SQL statement here:

```
"select customers.companyname, customers.contactname,
customers.city, products.productname from customers inner join
orders on customers.customerid = orders.customerid inner join
`order details` on orders.orderid = `order details`.orderid inner
join products on `order details`.productid = products.productid"
## [1] "select customers.companyname, customers.contactname,\n      customers.city, products.productname from c
```

```
library(RSQLite)
drv <- dbDriver("SQLite")
con <- dbConnect(drv, dbname="extdata/Northwind.sl3")
# Check all tables using
dbListTables(con)
## [1] "Alphabetical list of products" "Categories"
## [3] "Current Product List"        "Customer and Suppliers by City"
## [5] "CustomerCustomerDemo"       "CustomerDemographics"
## [7] "Customers"                   "EmployeeTerritories"
## [9] "Employees"                   "Order Details"
## [11] "Order Details Extended"      "Order Subtotals"
## [13] "Orders"                      "Orders Qry"
## [15] "Products"                    "Products Above Average Price"
## [17] "Products by Category"        "Region"
## [19] "Shippers"                    "Summary of Sales by Quarter"
## [21] "Summary of Sales by Year"     "Suppliers"
## [23] "Territories"                 "copy_of_customers"
# Check fields of a table
dbListFields(con, 'Customers')
## [1] "CustomerID" "CompanyName" "ContactName" "ContactTitle"
## [5] "Address"    "City"          "Region"      "PostalCode"
## [9] "Country"    "Phone"         "Fax"
dbListFields(con, 'Products')
## [1] "ProductID" "ProductName" "SupplierID" "CategoryID"
## [5] "QuantityPerUnit" "UnitPrice" "UnitsInStock" "UnitsOnOrder"
## [9] "ReorderLevel" "Discontinued"

tab <- dbGetQuery(con, "select customers.companyname, customers.contactname,
                        customers.city, products.productname from customers inner join
                        orders on customers.customerid = orders.customerid inner join
                        `order details` on orders.orderid = `order details`.orderid inner
```



## Data Analysis and Visualization Exercise 2.2: Data import

```
join products on `order details`.productid = products.productid")
nrow(tab)
## [1] 2155
tab[1:5,]
##           CompanyName  ContactName      City
## 1 Vins et alcools Chevalier  Paul Henriot    Reims
## 2 Vins et alcools Chevalier  Paul Henriot    Reims
## 3 Vins et alcools Chevalier  Paul Henriot    Reims
## 4 Toms Spezialit\xe4ten  Karin Josephs  M\xfcnster
## 5 Toms Spezialit\xe4ten  Karin Josephs  M\xfcnster
##           ProductName
## 1 Queso Cabrales
## 2 Singaporean Hokkien Fried Mee
## 3 Mozzarella di Giovanni
## 4 Tofu
## 5 Manjimup Dried Apples
```

## 4 XML questions

### 4.1 Question

Load the XML document `plant_catalog.xml`. Use XPath and DOM functions to find out all unique element names in the document.

```
library(XML)
doc = xmlTreeParse("extdata/plant_catalog.xml", useInternal = TRUE)
#root = xmlRoot(doc)
unique(unlist(xpathApply(doc, "//*", xmlName)))
## [1] "CATALOG"      "PLANT"        "COMMON"       "BOTANICAL"
## [5] "ZONE"        "LIGHT"        "PRICE"        "AVAILABILITY"
```

Get all plants of zone 4 and transform the data into an R list. Hint: `'xmlToList'`

```
plant_list = xpathApply(doc, "//PLANT[ZONE = 4]", xmlToList)
length(plant_list)
## [1] 15
plant_list[[1]]
## $COMMON
## [1] "Bloodroot"
##
## $BOTANICAL
## [1] "Sanguinaria canadensis"
##
## $ZONE
## [1] "4"
##
## $LIGHT
## [1] "Mostly Shady"
```

## Data Analysis and Visualization Exercise 2.2: Data import

```
##  
## $PRICE  
## [1] "$2.44"  
##  
## $AVAILABILITY  
## [1] "031599"
```

### 4.2 Question

Read the HTML tables from the website <https://www.skysports.com/premier-league-table> into your workspace.

Which team is currently placed first in the premier league?

```
library(XML)  
library(RCurl)  
url <- "https://www.skysports.com/premier-league-table"  
tables <- readHTMLTable(getURL(url))  
table <- tables[[1]]  
table[1,"Team"]  
## [1] Liverpool  
## 20 Levels: Arsenal Aston Villa Bournemouth ... Wolverhampton Wanderers
```

## 5 Prepare for next lecture

---

For data manipulation with the `data.table` package please read this [intro](#).